

Content-Addressed Continual Memory with On-Demand Capacity

Xinyu Yang
pana.yang@hotmail.com

August 29, 2026

Abstract

We study a continual-learning memory in which storage is addressed by content, capacity is allocated on demand from the stream without task identifiers, and new capacity is appended without perturbing anything already stored. On a relational stream whose regimes are distinguishable from their inputs, the mechanism retains 49.5% of probed facts zero-shot where a parameter-matched monolithic model retains 15.0%. Sharing the payload transform across regimes while keeping per-regime readouts adds 16.4 points over a fully private variant (six seeds, $p < 0.001$). Allocating readout rows only for targets a regime actually emits reduces occupied capacity by $7.6\times$ at no cost in accuracy, and appending capacity leaves existing weights byte-identical and existing predictions bit-identical.

We report the negative results at equal length. The architecture carries nothing between facts beyond a single-timescale context summary, and we quantify the resulting loss of association across a gap. Multi-hop routing, which could in principle purchase combinatorial addressing from linear memory, does not: it occupies the same capacity for 25 points less accuracy. Retention declines as regimes accumulate, and we show by measurement that the decline is neither capacity starvation, nor worsening misrouting, nor failing separation, but retrieval difficulty growing sub-proportionally with what is stored.

1 Introduction

A learner that must absorb an unbounded stream with bounded parameters has to discard something. Regularisation methods such as elastic weight consolidation [Kirkpatrick et al., 2017] give up plasticity, protecting old weights at the cost of new learning. Methods that simply keep training give up the past. A third option is to *allocate*: give distinct regimes distinct storage, so that degradation is uniform rather than selective. Dynamically expanding architectures [Rusu et al., 2016, Yoon et al., 2018] take this route and are routinely criticised for adding parameters per task without demonstrating that the addition is proportionate to new content rather than to task count.

This paper examines an allocation mechanism built around three commitments. Addressing is by content through a fixed random hash, so a fact reaches the same storage under every history. Regime identity is inferred online from the stream, so no task identifier, task boundary, or replay buffer is required. Capacity is appended rather than resized, so adding storage provably does not disturb what is already there.

We report what these commitments buy and what they cost. They buy a large advantage where regimes are distinguishable, a capacity footprint that follows emitted content rather than vocabulary size, and a per-token cost independent of sequence length. They cost the ability to associate items across a gap, which content-only addressing forbids by construction, and they leave retention declining as regimes accumulate.

Contributions.

1. An allocation mechanism whose two addressing factors — content to a memory slice, inferred regime to a readout — are separately measurable, with no task supervision at any point (§3).
2. Evidence that sharing the payload transform while privatising the readout is the correct split for a stream whose regimes share surface form and differ in targets, confirmed from both directions (§5.2).
3. A capacity result: sampling negatives from a regime’s own emitted targets makes occupancy follow content, $7.6\times$ smaller at equal accuracy (§5.3).
4. A non-destructiveness property, verified exactly rather than statistically: appending capacity leaves stored weights and predictions unchanged bit for bit (§5.4).
5. Measurements of what the mechanism does not currently do — cross-gap association, multi-hop addressing, and retention as regimes accumulate — reported with the alternative explanations we could exclude, and left open where we could not (§6).

2 Related work

Continual learning. Catastrophic interference in distributed representations [McCloskey and Cohen, 1989, French, 1999] motivates three broad families [Parisi et al., 2019, De Lange et al., 2021]: regularisation [Kirkpatrick et al., 2017], rehearsal, and architectural expansion [Rusu et al., 2016, Yoon et al., 2018]. Our mechanism is architectural, but differs from prior expansion methods in that it requires no task boundary to decide when to expand; the decision is made from the dispersion of an inferred regime’s own statistics. Task-free continual learning [Aljundi et al., 2019] shares this aim.

Conditional computation. Structurally the mechanism is a single-layer, top-1 sparsely-gated mixture of experts [Shazeer et al., 2017, Fedus et al., 2022] whose router is a fixed random hash rather than a learned gate, and whose expert count grows. We treat the unlearned router as a limitation to be reported rather than a feature: it forgoes load balancing and learned specialisation, and it is precisely what prevents context from influencing retrieval (§6).

Addressable and distributed memory. Content-addressed storage with sparse codes follows sparse distributed memory [Kanerva, 1988] and hyperdimensional computing [Kanerva, 2009]. Recoverable association of *pairs*, which a smoothed context summary cannot provide, is available through holographic reduced representations [Plate, 1995] or fast-weight outer products [Ba et al., 2016, Schlag et al., 2021]; we identify this as the missing ingredient rather than claim it.

Consolidation across timescales. Cascades of variables at multiple timescales [Fusi et al., 2005, Benna and Fusi, 2016] target the capacity–lifetime tradeoff of bounded synapses. We tested such a cascade in two roles, as weight consolidation and as input context, and report both outcomes.

Attention and its costs. Self-attention [Vaswani et al., 2017] pays $O(N)$ per token in context access and degrades on material in the middle of long contexts [Liu et al., 2024]. A content-addressed memory pays a cost independent of N ; we quantify the crossover in §5.5.

Vocabulary growth. That distinct types grow sublinearly in a Zipf-distributed stream is Heaps’ law [Heaps, 1978]; we use it to ask whether allocated capacity tracks content, and report that in our generator it does not, for reasons traced to the generator itself (§6).

3 Method

A fact is a pair (e, r) with target t . Token embeddings are fixed random unit vectors and are never learned. The payload is

$$p_0 = \text{normalise}(\text{emb}[e] + \text{emb}[r]). \quad (1)$$

3.1 Two addressing factors

Content to edge. Nodes of a small-world ring hold fixed random keys $k_{n,s}$. The walk selects $s^* = \arg \max_s \langle k_{n,s}, p_0 \rangle$ for H hops. Nothing in this path is learned, so a fact reaches the same edge under every history and routing consistency is 1.000 by construction.

Regime to readout. A slow exponential average c of stream tokens is matched against prototypes, $\gamma = \arg \max_j \langle \pi_j, c \rangle$, selecting a readout slice. No task identifier is used anywhere; the mechanism is never told which regime is live.

3.2 Memory

The edge transform is shared across regimes and applied residually,

$$p \leftarrow \text{normalise}(p + \tanh(E_{\text{edge}} p)), \quad (2)$$

while the readout R_γ is private per regime and produces logits $R_\gamma p_H$. The split follows the generative structure rather than preference: a payload is built from entity and relation, which are common across regimes, whereas targets are exactly what regimes disagree about. Section 5.2 confirms this from both directions.

3.3 Learning

Training is online, single-pass and prequential: predict, take the loss, then write. There is no replay buffer, no task boundary and no second pass. The readout takes a delta-rule step on the target row together with negatives sampled from targets that regime has already emitted; gradients propagate back along the walked path into the shared transform. Routers receive no gradient.

3.4 Allocation

A class is split when the dispersion of its own similarity distribution, $\sigma_\gamma/|\mu_\gamma|$, exceeds a threshold, and the new prototype is placed at the current context. Dispersion is used rather than an outlier test because outlier tests fail in two ways we measured: against a global running distribution they fail once regimes are numerous, since the distribution becomes their mixture; against a class’s own distribution they fail because a class absorbing more regimes raises its own variance, so the threshold rises exactly when splitting becomes necessary. Allocation appends, leaving existing slices at their offsets and values.

4 Experimental setup

The source is a multi-domain relational graph walked continuously, with facts stratified by visit frequency into hub, mid and tail tiers under a Zipf distribution. Two modes are used. In **Mode A** each regime owns its entities, so regime identity is present in the input. In **Mode B** entities and relations are byte-identical across regimes and only targets differ, which removes the routing

signal and, by construction, any shared target to amortise; we use it as an adversarial bound rather than as the main setting.

Evaluation is zero-shot retention of probed facts under a frozen memory, with the probe restoring the context it needs and returning the state it displaced. All figures in this paper are extracted by a single parser that anchors on the results table header and reads the column by name.

5 Results

5.1 Where regimes are distinguishable

regimes	this work	monolithic	advantage
8	49.53%	15.10%	+34.43pp
16	43.07%	12.50%	+30.57pp

Table 1: Mode A, three seeds at eight regimes and one at sixteen. The monolithic control sits at the level of a linear baseline without sparse coding (15.30%), confirming the task is not solvable by predicting frequent targets. Baselines at 32 and 64 regimes are still running.

5.2 Sharing the transform, privatising the readout

Mode B, twelve regimes, six seeds	zero-shot	paired difference
Shared transform, private readouts	29.42%	+16.43pp, 6/6, $p < 0.001$
Fully private	12.98%	—
Monolithic	21.28%	+8.13pp, 5/6, $p = 0.119$

Table 2: Sharing the transform is established against the private architecture. The advantage over the monolithic control is not.

Sharing the *readout* as well costs 27.50 points. A common readout block is updated on every write and therefore accumulates every regime’s target mapping, while Mode B is defined by the same (e, r) pointing at different targets per regime; it averages contradictions. The asymmetry of §3 is thus confirmed from both sides, one by making the change and the other by testing its opposite.

5.3 Capacity follows content

	occupied rows	rows per class	accuracy
Dense delta rule	45056/45056 (100%)	4096	32.6%
Sampled from emitted targets	5918/45056 (13.1%)	538	33.8%

Table 3: A $7.6\times$ reduction in occupied capacity at no cost in accuracy.

5.4 Expansion does not disturb what is stored

Appending a class slice leaves every existing weight byte-identical and **forward** bit-identical, asserted as a test rather than measured statistically. A monolithic model cannot offer this: its

capacity is the sparse-coding width, and widening it changes the projection through which every stored code was written.

5.5 Cost per token is independent of context length

Context access costs Hd^2 per token regardless of sequence length, against Nd for attention; the crossover is at $N = d$ and the ratio is $128\times$ at $N = 4096$. Against ourselves: overall activation is 6.178% of memory parameters, of which the edge memory contributes 1024 and the remainder is the readout scoring all V rows. Sparsity currently holds for the expert part only.

6 Open questions

Every negative reading in this study that we examined closely turned out, at least once, to be produced by the measurement rather than by the mechanism (Appendix A). We therefore report the following as measurements whose interpretation is unsettled, rather than as properties of the architecture.

Can the mechanism associate across a gap? Nothing is carried between facts beyond a single-timescale context, so a cue survives $(1 - \rho)^{\text{gap}}$.

gap	cue remaining	accuracy
1	0.990	10.10%
16	0.851	10.33%
64	0.526	8.40%
256	0.076	6.83%

Table 4: Paired degradation +3.27pp, 5/6 seeds, $p = 0.013$. The location of the cliff was predicted from the context rate before the run.

Exact association under constant memory is information-theoretically unavailable, so the design question is what a lossy code should retain. A multi-scale average answers *which regime* and cannot answer *which item*. Placing context into the address recovers roughly a tenth of the gap and costs lookup accuracy, consistent with content-only addressing being what forbids association in the first place.

Whether this is a limit of the mechanism or of this particular probe is not settled. Three earlier versions of the same measurement were invalid, each for a different reason, and the fourth is the one reported here; placing context into the address recovers roughly a tenth of the gap, which suggests the addressing scheme rather than the context code may be the binding constraint.

hops	zero-shot	occupied rows
1	55.15%	750
2	46.00%	737
3	29.63%	750

Table 5: Identical occupancy, monotonically worse retention. The under-training explanation is excluded: writes per edge *rise* with hops, from 3246 to 9222, because each fact touches more edges.

Does multi-hop addressing ever pay? We do not conclude that depth is unnecessary. Multi-hop is the only form of depth this architecture has, an earlier version of this comparison was taken while regime separation had collapsed, and the present one is on a single source at one scale. Whether combinatorial addressing pays under a source that rewards composition is untested.

regimes	zero-shot	occupied	rows/regime	home classes per regime
8	49.53%	530	66	0.75
16	43.07%	1142	71	0.64
32	37.65%	2412	75	0.70
64	34.80%	5198	81	0.78

Table 6: Monolithic baselines beyond eight regimes are in progress.

Why does retention decline as regimes accumulate? Three explanations are excluded by the data. Capacity is supplied linearly, so this is not starvation. Misrouted write magnitude is 57.0, 54.7, 51.0, 55.0 per cent — flat — so it is not worsening interference. Home classes per regime is flat, so separation does not degrade. What remains is retrieval: confusable stored targets grow $9.8\times$ while accuracy falls to $0.70\times$, and distinct predictions rise from 47 to 57 rather than collapsing.

These exclusions constrain the explanation without establishing it, and the monolithic baselines beyond sixteen regimes are still running. We state the retrieval account as the surviving candidate rather than as the finding.

Does acquiring a regime become cheaper as regimes accumulate? Later regimes do cost less per unit accuracy in our runs, but our measurement staggers arrival, and arrival order is itself sufficient to produce the effect: a regime arriving late enters a system whose transform and prototypes are already shaped. Separating order from content requires a design we have not run, so we leave the question open rather than answer it from a measurement that cannot distinguish the two (Appendix B).

Is the decline with scale a power law? Retention falls roughly as $N^{-\alpha}$ with $\alpha \approx 0.16$ over the measured range, N being the number of stored items. Whether this exponent is a property of the mechanism, and whether stronger addressing rather than more capacity moves it, is measurable without new machinery.

What should a constant-memory context retain? Exact association across a gap is information-theoretically unavailable under constant memory, so the design question is what to keep. A multi-scale average answers *which regime* and cannot answer *which item*. Sparse retention of salient items, or superposition with approximate unbinding [Plate, 1995, Ba et al., 2016], keeps identity where an average cannot.

Does the plasticity–stability frontier extend or shift? The frontier traced by our mechanism and that traced by the monolithic control barely overlap in plasticity, so at present only reachability can be compared, not position.

7 Limitations

Results are on a synthetic relational source; no natural-language corpus is evaluated. The monolithic baseline is established at eight regimes and is pending at larger scales. Cost comparisons

between the dense and sampled write paths are not commensurable, because write magnitude accumulates over a different number of rows in each. The private and shared arms differ in parameter count, which strengthens the accuracy comparison but confounds comparisons of cost. Whether acquiring a regime becomes cheaper as regimes accumulate is not settled here: our measurement staggers arrival, and arrival order alone can produce the effect (Appendix B).

8 Conclusion

Content addressing with on-demand allocation buys a large advantage where regimes are distinguishable, a capacity footprint that follows emitted content, and a per-token cost independent of context length, while provably not disturbing what is already stored. The same commitment that makes storage stable — addressing by content alone — is what prevents association across a gap, and the cost of that trade is measured here rather than asserted.

A Measurements invalidated during this study

Six measurements were invalidated after collection and re-run. Three were probe-side leaks in which a per-domain snapshot keyed by ground truth handed the mechanism the regime identity it was meant to infer. One was the allocation criterion of §3.4 in an earlier form, which compared novelty against a global running distribution and collapsed every regime into a single class at scale, voiding the savings, capacity and multi-hop measurements taken under it. One was a parser that scanned past the results table into a table with the same column count and reported 54.2% as 100.00%.

Each was caught by a contradiction between statistics rather than by reading code: a count that disagreed with another count, a diagnostic that printed nothing, an accuracy that could not coexist with its own exposure figures. The instruments that survived — write locality by time since regime change, occupancy against emitted content, allocation timing against regime boundaries — are reported alongside the results because they are what made the failures visible.

B Arrival order and the cost of a late regime

Measuring whether a later regime is cheaper to acquire requires staggered arrival, and staggering introduces a confound we record here for the benefit of anyone running the same measurement.

	shared transform	private
Mode A, staggered arrival	0.754	0.820
Mode B, disjoint targets, staggered arrival	0.442	0.505
Mode A, no staggering (index only)	1.199	1.207

Table 7: Cost of the late regimes relative to the early ones; below one means a late regime is cheaper. Three seeds per control.

Mode B’s targets are disjoint across regimes by construction, so nothing there can be amortised, yet the ratio is lower than in Mode A. Removing the staggering removes the effect. A regime arriving late therefore benefits from entering an already-shaped system, independently of whether it shares content with its predecessors, and a ratio below one cannot on its own be read as amortisation.

References

- R. Aljundi, K. Kelchtermans, and T. Tuytelaars. Task-free continual learning. In *CVPR*, 2019.
- J. Ba, G. Hinton, V. Mnih, J. Z. Leibo, and C. Ionescu. Using fast weights to attend to the recent past. In *NeurIPS*, 2016.
- M. K. Benna and S. Fusi. Computational principles of synaptic memory consolidation. *Nature Neuroscience*, 2016.
- M. De Lange, R. Aljundi, M. Masana, S. Parisot, X. Jia, A. Leonardis, G. Slabaugh, and T. Tuytelaars. A continual learning survey: Defying forgetting in classification tasks. *IEEE TPAMI*, 2021.
- W. Fedus, B. Zoph, and N. Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *JMLR*, 2022.
- R. M. French. Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 1999.
- S. Fusi, P. J. Drew, and L. F. Abbott. Cascade models of synaptically stored memories. *Neuron*, 2005.
- H. S. Heaps. *Information Retrieval: Computational and Theoretical Aspects*. Academic Press, 1978.
- P. Kanerva. *Sparse Distributed Memory*. MIT Press, 1988.
- P. Kanerva. Hyperdimensional computing: An introduction to computing in distributed representation with high-dimensional random vectors. *Cognitive Computation*, 2009.
- J. Kirkpatrick, R. Pascanu, N. Rabinowitz, J. Veness, G. Desjardins, A. A. Rusu, K. Milan, J. Quan, T. Ramalho, A. Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *PNAS*, 2017.
- N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: How language models use long contexts. *TACL*, 2024.
- M. McCloskey and N. J. Cohen. Catastrophic interference in connectionist networks: The sequential learning problem. *Psychology of Learning and Motivation*, 1989.
- G. I. Parisi, R. Kemker, J. L. Part, C. Kanan, and S. Wermter. Continual lifelong learning with neural networks: A review. *Neural Networks*, 2019.
- T. A. Plate. Holographic reduced representations. *IEEE Transactions on Neural Networks*, 1995.
- A. A. Rusu, N. C. Rabinowitz, G. Desjardins, H. Soyer, J. Kirkpatrick, K. Kavukcuoglu, R. Pascanu, and R. Hadsell. Progressive neural networks. *arXiv:1606.04671*, 2016.
- I. Schlag, K. Irie, and J. Schmidhuber. Linear transformers are secretly fast weight programmers. In *ICML*, 2021.
- N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *ICLR*, 2017.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *NeurIPS*, 2017.
- J. Yoon, E. Yang, J. Lee, and S. J. Hwang. Lifelong learning with dynamically expandable networks. In *ICLR*, 2018.